



ROI TRAINING
MAXIMIZE YOUR TRAINING INVESTMENT

Python Program

CHAPTER 3: GETTING STARTED WITH PANDAS

Chapter Objectives

In this chapter, we will introduce:

- ➡ Pandas data structures
- ➡ Essential functionality of pandas
- ➡ Reading from data sources
- ➡ Summarizing and computing descriptive statistics
- ➡ Handling missing data

Chapter Concepts

Introduction to Pandas Data

Essential Functionality

Reading From Data Sources

Summarizing and Computing Descriptive Statistics

Handling Missing Data

Chapter Summary

Introducing Pandas

- ➔ *Pandas* is an open-source, BSD-licensed library
 - Provides high-performance, easy-to-use data structures and data analysis tools
- ➔ Common usages
 - `import pandas as pd`
 - `from pandas import DataFrame`
 - `from pandas import Series`
 - `from pandas import DataFrame, Series`

Pandas Series Data Structure

- ➔ Pandas provides the `Series` data structure
 - A one-dimensional array-like object
- ➔ Contains:
 - Array of data
 - Array of data labels known as the index
- ➔ `Series` object has `values` and `index` properties
- ➔ Can be thought of as a fixed-length dictionary

Series Example

```
from pandas import Series
```

```
data = Series([1,2,3,4])
```

```
print(data)
```

```
0    1
```

```
1    2
```

```
2    3
```

```
3    4
```

```
dtype: int64
```

Each data value is assigned an index from N through to N-1

```
print(data.values)
```

```
array([1, 2, 3, 4])
```

```
print(data.index)
```

```
RangeIndex(start=0, stop=4, step=1)
```

```
print(data[1])
```

```
2
```

Series Index

- ➔ It is possible to create a `Series` with a user-defined index for each data point

```
data = Series([1,2,3,4], index=['a','b','c','d'])  
print(data.index)  
Index([u'a', u'b', u'c', u'd'], dtype='object')  
  
print(data['c'])  
3
```

Index can be used to access data points

Series and Dictionaries

- ➔ A `Series` can be created by passing a dictionary
 - Dictionary keys are used for `Series` index by default
 - ➔ Separate keys can be provided

```
cities = {'Dublin': 200000, 'Athlone': 15000, 'Galway': 700000}
series1 = Series(cities)
print (series1)
Athlone      15000
Dublin       200000
Galway       700000
dtype: int64
```


Series and Dictionaries (continued)

- ➔ If no key is provided, then NaN is used
 - Can use `isnull()` and `notnull()` functions to detect missing data

```
cities = {'Dublin': 200000, 'Athlone': 15000, 'Galway': 700000}
```

```
indexes = ['Dublin', 'Athlone', 'Waterford']
```

```
series2 = Series(cities, index=indexes)
```

```
print (series2)
```

```
Dublin          200000.0
Athlone          15000.0
Waterford       NaN
dtype: float64
```

Missing value

Detecting Missing Data

```
print(series2)
Dublin          200000.0
Athlone         15000.0
Waterford              NaN
dtype: float64

print(series2.isnull())
Dublin          False
Athlone         False
Waterford       True
dtype: bool

print(series2.notnull())
Dublin          True
Athlone         True
Waterford       False
dtype: bool
```

DataFrame

- ➡ DataFrame represents a tabular data structure
 - Similar to spreadsheet
 - Contains ordered collection of rows and columns
- ➡ Has both a row and column index
- ➡ Most common way to construct is from a dictionary of lists or NumPy arrays
 - Must be equal length
 - Index will be provided automatically
 - Columns placed in sorted order by default

DataFrame Example

```
from pandas import DataFrame

data = {'team':['Leicester', 'Manchester City',
'Arsenal'], 'player':['Vardy', 'Aguero', 'Sanchez'],
'goals':[24,22,19]}

football = DataFrame(data)

print(football)
```

	goals	player	team
0	24	Vardy	Leicester
1	22	Aguero	Manchester City
2	19	Sanchez	Arsenal

DataFrame Indexes

➔ Column order can be specified when creating DataFrame

```
data = {'team': ['Leicester', 'Manchester City',  
               'Arsenal'], 'player': ['Vardy', 'Aguero', 'Sanchez'],  
        'goals': [24, 22, 19]}
```

```
football = DataFrame(data,  
                     columns=['player', 'team', 'goals', 'played'],  
                     index=['one', 'two', 'three'])
```

```
print(football)
```

	player	team	goals	played
one	Vardy	Leicester	24	NaN
two	Aguero	Manchester City	22	NaN
three	Sanchez	Arsenal	19	NaN

Extra column

Extra column

Index Objects

- Index objects are immutable and cannot be modified
 - Can be shared across data structures
 - Act as a set

```
print(football)
```

	player	team	goals	played
one	Vardy	Leicester	24	NaN
two	Aguero	Manchester City	22	NaN
three	Sanchez	Arsenal	19	NaN

```
print('player' in football.columns)
```

```
True
```

Set operations

```
print('three' in football.index)
```

```
True
```

- Index has a number of methods found at:
 - <http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Index.html>

Chapter Concepts

Introduction to Pandas Data

Essential Functionality

Reading From Data Sources

Summarizing and Computing Descriptive Statistics

Handling Missing Data

Chapter Summary

Series Indexing and Selection

➔ Data from `Series` can be retrieved by index using integers and indexes

```
data = Series(np.arange(4.0), index=['a', 'b', 'c', 'd'])
```

```
print(data)
a      0.0
b      1.0
c      2.0
d      3.0
dtype: float64
```

```
print(data[2])
2.0
```

```
print(data[['b', 'd']])
b      1.0
d      3.0
dtype: float64
```

```
print(data<2)
a      True
b      True
c     False
d     False
```

```
print(data[data<2])
a      0.0
b      1.0
dtype: float64
```


DataFrame Indexing and Selection

➔ Indexing retrieves one or more columns

```
data = DataFrame(np.arange(9).reshape((3,3)),  
                  index=['a','b','c'], columns=['one','two','three'])
```

```
print(data)
```

	one	two	three
a	0	1	2
b	3	4	5
c	6	7	8

```
print(data['three'])
```

a	2
b	5
c	8

Name: three, dtype:
int64

```
print(data[:2])
```

	one	two	three
a	0	1	2
b	3	4	5

```
print(data['two']>1)
```

a	False
b	True
c	True

Name: two, dtype:bool

```
print(data[data['two']>1])
```

	one	two	three
b	3	4	5
c	6	7	8

DataFrame Indexing with ix, loc, and iloc

- ➔ DataFrame has several indexing fields: ix, loc, and iloc
 - Allows selecting subset of rows and columns

```
data = DataFrame(np.arange(9).reshape((3,3)),  
                 index=['a','b','c'], columns=['one','two','three'])
```

	one	two	three
a	0	1	2
b	3	4	5
c	6	7	8

- ➔ All these could be used to retrieve the first row
 - data.ix[0] data.ix['a'] data.loc['a'] data.iloc[0]
- ➔ All these would retrieve the second column for all rows
 - data.ix[:, 'two'], data.ix[:, 1], data.loc[:, 'two'], data.iloc[:, 1]
- ➔ All these would retrieve the first two rows and the second column
 - data.ix[['a','b'], 'two'], data.ix[0:2, 1], data.loc['a':'c', 'two'], data.iloc[0:2, 1]

Series

Arithmetic and Data Alignment

- ➔ When adding together objects if the index pairs are not the same, then index in result is the union of the index pairs

```
data1 = Series([1.0,2.0,3.0], index=['a','d','e'])  
data2 = Series([2.0,3.0,4.0, 5.0], index=['a','b','c','e'])
```

```
print(data1)  
a      1.0  
d      2.0  
e      3.0  
dtype: float64
```

```
print(data2)  
a      2.0  
b      3.0  
c      4.0  
e      5.0  
dtype: float64
```

```
print(data1 + data2)  
a      3.0  
b      NaN  
c      NaN  
d      NaN  
e      8.0  
dtype: float64
```

NaN for indices that
do not overlap

DataFrame

Arithmetic and Data Alignment

➔ DataFrame's alignment is performed on columns and rows

```
data1 = DataFrame(np.arange(9.0).reshape((3,3)),  
                  columns=list('abc'), index=['one', 'two', 'three'])
```

```
data2 = DataFrame(np.arange(12.0).reshape((4,3)),  
                  columns=list('ace'), index=['one', 'two', 'three', 'four'])
```

```
print(data1)
```

	a	b	c
one	0.0	1.0	2.0
two	3.0	4.0	5.0
three	6.0	7.0	8.0

```
print(data2)
```

	a	c	e
one	0.0	1.0	2.0
two	3.0	4.0	5.0
three	6.0	7.0	8.0
four	9.0	10.0	11.0

```
print(data1 + data2)
```

	a	b	c	e
four	NaN	NaN	NaN	NaN
one	0.0	NaN	3.0	NaN
three	12.0	NaN	15.0	NaN
two	6.0	NaN	9.0	NaN

Index and columns
are unions of data1
and data2

Arithmetic with Fill Values

- ➔ For arithmetic between differently indexed objects, can use `fill_value` to prevent missing values appearing in resultant data structure

```
print(data1)
```

	a	b	c
one	0.0	1.0	2.0
two	3.0	4.0	5.0
three	6.0	7.0	8.0

```
print(data2)
```

	a	c	d	e
one	0.0	1.0	2.0	3.0
two	4.0	5.0	6.0	7.0
three	8.0	9.0	10.0	11.0

```
data1.add(data2, fill_value=0)
```

	a	b	c	d	e
one	0.0	1.0	3.0	2.0	3.0
two	7.0	4.0	10.0	6.0	7.0
three	14.0	7.0	17.0	10.0	11.0

Function Application and Mapping

- ➔ A frequent operation is to apply a function to each column or row of DataFrame

```
data = DataFrame(np.random.randn(4,4))
print(data)
```

	0	1	2	3
0	0.546781	-0.862394	-2.384923	-0.098065
1	-0.219738	0.172776	-1.558335	-0.124880
2	-1.016070	-0.670825	-1.602997	-0.018526
3	-0.050491	0.258218	0.073574	0.144686

```
f = lambda x: x.max() - x.min()
```

```
print(data.apply(f))
print(data.apply(f, axis=0))
```

0	1.562851
1	1.120612
2	2.458497
3	0.269566

dtype: float64

Apply per
column

```
print(data.apply(f, axis=1))
```

0	2.931705
1	1.731110
2	1.584471
3	0.308709

dtype: float64

Along row

Sorting

- ➔ Sorting by index on either axis is available, ascending order by default

```
print(data)
```

	0	1	2	3
0	0.546781	-0.862394	-2.384923	-0.098065
1	-0.219738	0.172776	-1.558335	-0.124880
2	-1.016070	-0.670825	-1.602997	-0.018526
3	-0.050491	0.258218	0.073574	0.144686

Specify axis = 1 for
column sorting

```
print(data.sort_index(ascending=False))
```

	0	1	2	3
3	-0.050491	0.258218	0.073574	0.144686
2	-1.016070	-0.670825	-1.602997	-0.018526
1	-0.219738	0.172776	-1.558335	-0.124880
0	0.546781	-0.862394	-2.384923	-0.098065

Descending
order on rows

Could supply two sort
keys with a list

- ➔ Can also sort by values instead of index

```
print(data.sort_values(by=[1], ascending=False))
```

Ranking

- ➔ Ranking assigns ranks from 1 to the number of valid data points in an array

```
data = DataFrame({'b': [1,4,3,2], 'a': [6,9,20,3], 'c': [7,2,8,15]})  
print(data)
```

	a	b	c
0	6	1	7
1	9	4	2
2	20	3	8
3	3	2	15

Rank on column

Rank on row

```
print(data.rank())
```

	a	b	c
0	2.0	1.0	2.0
1	3.0	4.0	1.0
2	4.0	3.0	3.0
3	1.0	2.0	4.0

```
print(data.rank(axis=1))
```

	a	b	c
0	2.0	1.0	3.0
1	3.0	2.0	1.0
2	3.0	1.0	2.0
3	2.0	1.0	3.0

Chapter Concepts

Introduction to Pandas Data

Essential Functionality

Reading From Data Sources

Summarizing and Computing Descriptive Statistics

Handling Missing Data

Chapter Summary

Data

- To use the tools in this course, we need data to work with
- Data can be used from a variety of sources:
 - Text format
 - CSV
 - JSON
 - XML/HTML
 - Binary formats
 - HDF5
 - Excel
 - Databases
 - MongoDB

Working with Text Formats

- We will introduce pandas during this course
 - It provides a number of features for working with data in a tabular format
 - Known as a `DataFrame`
 - We will use this in our examples here with details to follow
- Pandas provides the following functions for reading data:
 - `read_csv`
 - Load data from delimited file or URL, comma is default delimiter
 - `read_table`
 - Load delimited data from a file or URL, tab is default delimiter
 - `read_fwf`
 - Read fixed-width column formatted file or URL

Reading a Text File Source

➔ The following code loads a `csv` file

```
data = pd.read_csv('sample.csv')  
print(data)
```

	1	2	3	4	hello
0	5	6	7	8	world
1	9	10	11	12	some message

Returns a DataFrame

➔ `read_table` would work for this file too, but it's been deprecated

```
pd.read_table('sample.csv', sep=',')
```

	1	2	3	4	hello
0	5	6	7	8	world
1	9	10	11	12	some message

Delimiter to use

Reading Large Files

- It is possible to read large files in smaller fragments
 - Specify a `chunksize` to `read_csv`
 - Size is number of lines to supply

```
fragment = pd.read_csv('sample.csv', chunksize=1)
```

```
for line in fragment:  
    print(line)
```



Read a line at a time

```
1 2 3 4 hello  
0 5 6 7 8 world  
1 2 3 4 hello  
0 9 10 11 12 some message
```

Writing Data Out to Text Files

- ➔ Data can be exported to files in a delimited format

```
print(data)
```

```
      1      2      3      4      hello  
0  5      6      7      8      world  
1  9     10     11     12  some message
```

```
data.to_csv('file1.csv')
```

Write to file

- ➔ Can specify a separator too

```
data.to_csv('file1.csv', sep = '|')
```

Separator parameter

JSON Data

- ➔ Can read in JSON data using Python
 - Create DataFrame from data
- ➔ Consider the following JSON file:

```
{ "name": "jayne",  
  "role": "sales",  
  "customers" :  
    [{"name": "Andersons", "product": "Bosch", "quantity": 100},  
     {"name": "ElectricalDirect", "product": "Miele",  
      "quantity": 200}]}
```

```
import json  
data = json.loads(open('example.json').read())  
customers = DataFrame(data['customers'])  
print(customers)
```

	name	product	quantity
0	Andersons	Bosch	100
1	Electrical Direct	Miele	200

Read JSON file

SQL Data

- ➔ Can read from SQL databases using a standard recipe
 - Import the package for the particular database
 - Open a connection
 - Create a cursor
 - Execute a query
 - Iterate through the cursor or fetch the results to a list
 - Close the connection when done

```
import sqlite3
cn = sqlite3.connect('test.sqlite')
curs = cn.cursor()
curs.execute("create table names (id int, name varchar(20))")
curs.execute("insert into names values(1, 'Alice'), (2, 'Bob')")
cn.commit()
curs.execute("select * from names")
names = curs.fetchall()
print(names)
names2 = pd.read_sql_query("select * from names", cn)
print(names2)
cn.close()
```

```
[(1, 'Alice'), (2, 'Bob')]
```

	id	name
0	1	Alice
1	2	Bob

Other Formats

- Python has many libraries for reading and writing different formats/sources
 - HTML and XML
- We will not cover the details here, but at a high level data can be processed from:
 - HTML/XML sources
 - Binary formats
 - Microsoft Excel files
 - Web APIs (JSON)
 - Databases
 - Relational
 - MongoDB

Chapter Concepts

Introduction to Pandas Data

Essential Functionality

Reading From Data Sources

Summarizing and Computing Descriptive Statistics

Handling Missing Data

Chapter Summary

Summarizing and Computing Descriptive Statistics

- ➔ Pandas objects have a set of common mathematical and statistical methods
- ➔ Most are reductions or summary statistics
 - Extract a single value, e.g., `sum()`
 - They exclude missing data

```
data = DataFrame([[1,np.nan],[3,4],[5,np.nan]],  
                  columns=['a','b'])
```

```
print(data)
```

	a	b
0	1	NaN
1	3	4.0
2	5	NaN

Sums columns

```
print(data.sum())
```

```
a      9.0  
b      4.0  
dtype: float64
```

Sums rows

```
print(data.sum(axis=1))
```

```
0      1.0  
1      7.0  
2      5.0  
dtype: float64
```

Pandas Mathematical Methods

- ➔ A few methods return multiple values, e.g., `describe()`

```
print (data)
```

	a	b
0	1	NaN
1	3	4.0
2	5	NaN

```
print (data.describe())
```

	a	b
count	3.0	1.0
mean	3.0	4.0
std	2.0	NaN
min	1.0	4.0
25%	2.0	NaN
50%	3.0	NaN
75%	4.0	NaN
max	5.0	4.0

Produces
summary statistics

- ➔ Full list of DataFrame methods found at:

– <http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.html>

Correlation and Covariance

- ➔ Correlation and covariance are computed from pairs of arguments
- ➔ Consider fetching data from Yahoo! Finance

```
import pandas_datareader.data as web
all_data = {ticker: web.get_data_yahoo(ticker) for ticker in
['AAPL', 'IBM', 'MSFT', 'GOOG']}
print(all_data['AAPL'])
```

Data returned

[2418 rows x 6 columns],
'AAPL':

Date	Open	High	Low	Close	Volume
2010-01-04	626.951088	629.511067	624.241073	626.751061	3927000
2010-01-05	627.181073	627.841071	621.541045	623.991055	6031900
....					
Date	Adj Close				
2010-01-04	313.062468				
2010-01-05	311.683844				

Percentage Price Change

➔ Consider calculating the percentage change in the daily price of the stocks

```
price = DataFrame({ticker:data['Adj Close'] for  
                    ticker, data in all_data.items()})
```

```
print(price)
```

Date	AAPL	GOOG	IBM	MSFT
2010-01-04	27.727039	313.062468	111.405000	25.555485
2010-01-05	27.774976	311.683844	110.059232	25.563741
2010-01-06	27.333178	303.826685	109.344283	25.406859

```
returns = price.pct_change()
```

Daily price change

```
print(returns.tail())
```

	AAPL	GOOG	IBM	MSFT
Date				
2017-03-23	-0.003536	-0.014477	0.000229	-0.002460
2017-03-24	-0.001987	-0.003853	-0.005663	0.001696
2017-03-27	0.001707	0.006238	-0.000345	0.001847

Correlation and Covariance

- ➔ DataFrame's `corr()` and `cov()` methods return a correlation or covariance matrix as a DataFrame

```
print(returns.corr())
```

	AAPL	GOOG	IBM	MSFT
AAPL	1.000000	0.409814	0.382086	0.389641
GOOG	0.409814	1.000000	0.402671	0.471145
IBM	0.382086	0.402671	1.000000	0.495369
MSFT	0.389641	0.471145	0.495369	1.000000

```
print(returns.cov())
```

	AAPL	GOOG	IBM	MSFT
AAPL	0.000267	0.000104	0.000075	0.000092
GOOG	0.000104	0.000242	0.000075	0.000106
IBM	0.000075	0.000075	0.000143	0.000085
MSFT	0.000092	0.000106	0.000085	0.000208

Chapter Concepts

Introduction to Pandas Data

Essential Functionality

Reading From Data Sources

Summarizing and Computing Descriptive Statistics

Handling Missing Data

Chapter Summary

Handling Missing Data

- ➔ Missing data is common in most data analysis applications
- ➔ Pandas tries to make working with missing data as painless as possible
 - The NaN (NA) value is used to represent missing data
- ➔ Two approaches to working with missing data:
 - Filter out missing data
 - Fill in missing data

Filtering Out Missing Data: Series

```
from numpy import nan as NA
data = Series([1,NA,2,3,4,NA])
print(data)
0      1.0
1      NaN
2      2.0
3      3.0
4      4.0
5      NaN
dtype: float64

print(data.dropna())
0      1.0
2      2.0
3      3.0
4      4.0
dtype: float64

data = data.dropna()
data.dropna(inplace=True)
```

Returns only non-null values and index values but does not modify the original data set

Reassign it back to the same variable to modify it or use `inplace=True`

Filtering Out Missing Data: DataFrame

➤ Useful parameters include:

- `axis` which defaults to 0 for rows or 1 for columns
- `how` which drops the row or column if any one value is NA or all are NA

```
data = DataFrame  
([ [1,2,3], [NA,5,NA], [NA,NA,NA],  
  [10,11,12] ])
```

```
print(data.dropna(how='all'))
```

	0	1	2
0	1.0	2.0	3.0
1	NaN	5.0	NaN
3	10.0	11.0	12.0

```
print(data.dropna(how='any'))
```

	0	1	2
0	1.0	2.0	3.0
3	10.0	11.0	12.0

```
data = DataFrame  
([ [1,2,NA], [NA,5,NA], [NA,12,NA],  
  [10,11,NA] ])
```

```
print(data.dropna(how='all',axis=1))
```

	0	1
0	1.0	2.0
1	NaN	5.0
2	NaN	12.0
3	10.0	11.0

```
print(data.dropna(how='any',axis=1))
```

	1
0	2.0
1	5.0
2	12.0
3	11.0

Filling In Missing Data

```
print(data)
```

	0	1	2
0	1.0	2.0	3.0
1	NaN	5.0	NaN
2	NaN	NaN	NaN
3	10.0	11.0	12.0

```
filled = data.fillna(0)  
print(filled)
```

	0	1	2
0	1.0	2.0	3.0
1	0.0	5.0	0.0
2	0.0	0.0	0.0
3	10.0	11.0	12.0

```
filled = data.fillna({0:10, 1:11, 2:12})  
print(filled)
```

	0	1	2
0	1.0	2.0	3.0
1	10.0	5.0	12.0
2	10.0	11.0	12.0
3	10.0	11.0	12.0

Supply dictionary with
column:value pair for different
fill values per column

Chapter Concepts

Introduction to Pandas Data

Essential Functionality

Reading From Data Sources

Summarizing and Computing Descriptive Statistics

Handling Missing Data

Chapter Summary

Chapter Summary

In this chapter, we have introduced:

- Pandas data structures
- Essential functionality of pandas
- Reading from data sources
- Summarizing and computing descriptive statistics
- Handling missing data